



COMP 346 – FALL 2019

DEADLOCK

1

TOPICS


- Deadlock brief
 - Simplest Example
 - Necessary Deadlock Conditions
 - Resource Allocation Graph
 - Deadlock Handling
- Deadlock with Semaphores
- Examples



DEADLOCK

- Simplest example:

- Two processes require two resources to complete (and release the resources).
- There are only two instances of these resources.
- If they acquire one resource each, they block indefinitely waiting for each other to release the other resource =>
 - **DEADLOCK**, one of the biggest problems in multiprogramming.



NECESSARY **DEADLOCK** CONDITIONS

- Recall which conditions must hold :
 - **Mutual Exclusion:** at least one process exclusively uses a resource.
 - **Hold and wait:** a process possesses at least one resources and requires more, which are held by others.
 - **No preemption:** resources are released only in voluntary manner by processes holding them.
 - **Circular wait:** $P_1 \rightarrow P_2 \rightarrow P_3 \rightarrow \dots \rightarrow P_N \rightarrow P_1$.




RESOURCE-ALLOCATION GRAPH

- An easy way to illustrate resource allocation and visually detect the deadlock situation(s)

- Example:

- $N+1$ resources
- N processes
- Every process needs 2 resources
- Upon acquiring 2 resources, a process releases them
- Is there a deadlock?



HANDLING DEADLOCKS

- Deadlock Ignorance
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection and Recovery



DEADLOCK IGNORANCE

- System designers pretend that deadlocks don't happen.
- If they do happen, they don't happen very often.
- Most common approach because it's cheap (in terms of human labor, complexity of the system and system's performance).
- Highest resource utilization.
- System administrator can simply kill deadlocked processes or restart the system if it's not responding.



DEADLOCK PREVENTION

- Recall necessarily deadlock conditions.
- Deadlock prevention algorithms **assure at least one of these conditions do not hold**, thus *preventing* occurrence of the deadlock.
- Possible Disadvantages:
 - Low device utilization
 - Reduced system's throughput



DEADLOCK AVOIDANCE

- Unlike deadlock prevention, deadlock avoidance algorithms do not watch for necessary deadlock conditions, but **use additional priority information about future resource requests**.
- This information will help the system to avoid entering the *unsafe state*.
- Disadvantages: again, lower resource utilization (because resources may **not be granted** sometimes even if they are unused).



DEADLOCK DETECTION AND RECOVERY

- Instead of preventing or avoiding deadlocks we **allow them to happen** and also provide **mechanisms to recover**.
- Problems:
 - How **often** do we run detection algorithms?
 - How **do** we recover?
 - What is the **cost** of detection and recovery?



1- DEADLOCK WITH SEMAPHORES

- Semaphore alpha = new Semaphore(1);
- Semaphore beta = new Semaphore(1);

Thread 1:

alpha.Wait();

beta.Wait();

DoSomething();

beta.Signal();

alpha.Signal();

Thread 2:


beta.Wait();

alpha.Wait();

DoSomething();

alpha.Signal();

beta.Signal();



A SOLUTION

```
Semaphore alpha = new Semaphore(1);  
Semaphore beta  = new Semaphore(1);
```

Thread 1:

```
alpha.Wait();
```

```
beta.Wait();
```

```
DoSomething();
```

```
beta.Signal();
```

```
alpha.Signal();
```

Thread 2:

```
alpha.Wait();
```

```
beta.Wait();
```

```
DoSomething();
```

```
beta.Signal();
```

```
alpha.Signal();
```

Example of Banker's Algorithm

5 processes P_0 through P_4 ;

3 resource types:

A (10 instances), B (5 instances), and C (7 instances).

Snapshot at time T_0 :

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	$A\ B\ C$	$A\ B\ C$	$A\ B\ C$
P_0	0 1 0	7 5 3	3 3 2
P_1	2 0 0	3 2 2	
P_2	3 0 2	9 0 2	
P_3	2 1 1	2 2 2	
P_4	0 0 2	4 3 3	

Example (Cont.)

The content of the matrix *Need* is defined to be *Max - Allocation*.

	<u>Need</u>
	A B C
P_0	7 4 3
P_1	1 2 2
P_2	6 0 0
P_3	0 1 1
P_4	4 3 1

Available = Available + Allocation

The system is in a safe state since the sequence $\langle P_1, P_3, P_4, P_2, P_0 \rangle$ satisfies safety criteria.



Example: P_1 Request (1,0,2)

Check that Request \leq Available (that is, $(1,0,2) \leq (3,3,2) \Rightarrow$ true.

	<u>Allocation</u>	<u>Need</u>	<u>Available</u>
	A B C	A B C	A B C
P_0	0 1 0	7 4 3	2 3 0
P_1	3 0 2	0 2 0	
P_2	3 0 1	6 0 0	
P_3	2 1 1	0 1 1	
P_4	0 0 2	4 3 1	

Executing safety algorithm shows that sequence $\langle P_1, P_3, P_4, P_0, P_2 \rangle$ satisfies safety requirement.

Can request for (3,3,0) by P_4 be granted?

Can request for (0,2,0) by P_0 be granted?



PROBLEM 1


- In operating systems that implement multi-threading, the resources are allocated to the controlling process but it is the threads that can become deadlocked.
 - (i) What is the advantage to allocate the resources to the process?
 - (ii) What is the advantage that only the threads can be deadlocked?

SOLUTION 1

- *(i) The resources can be shared by all the threads of a process.*
- *(ii) Blocking of a thread does not block the other threads.*

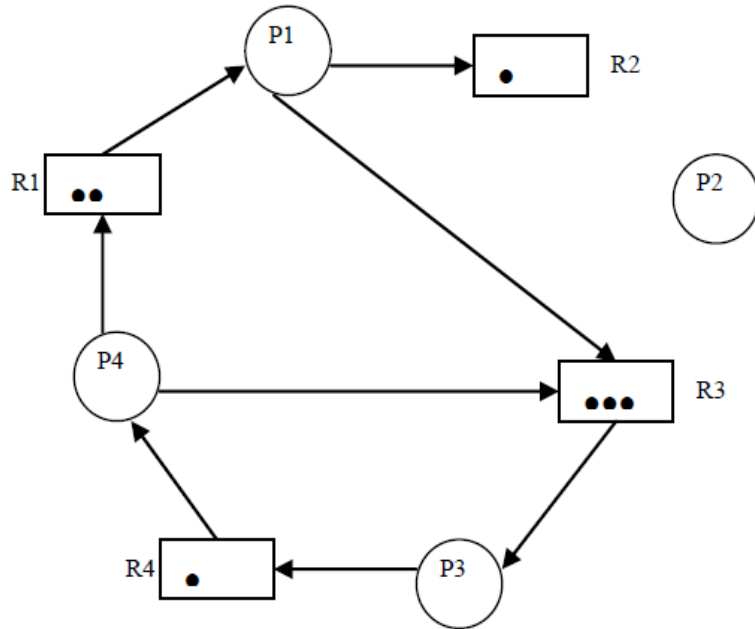
PROBLEM 2

- Consider the following resource allocation graph. Using the graph reduction method, show that there is no deadlock.

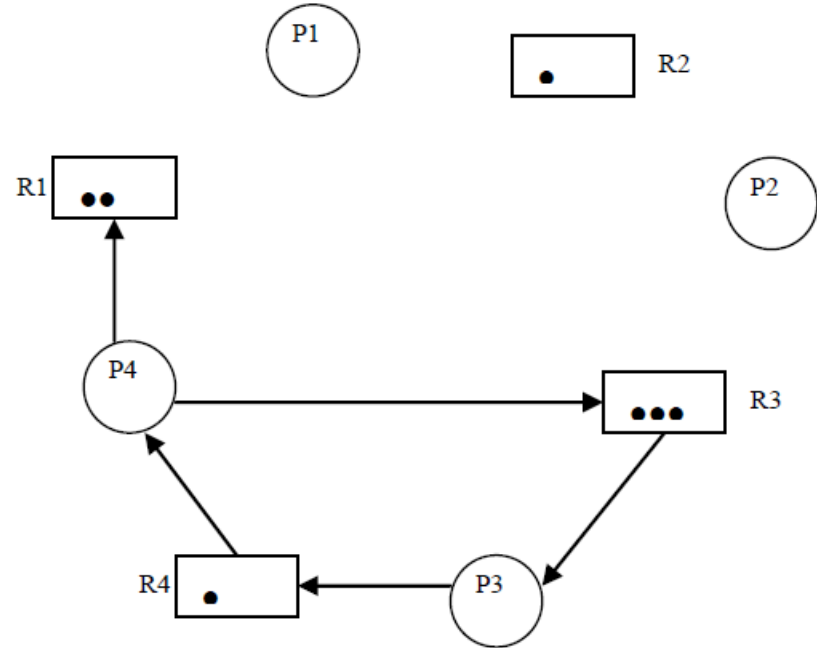


SOLUTION 2

- Step 1 : reduce the graph by P2.




- Step 2 : reduce the graph by P1.




SOLUTION 2 - CONTINUE

- Step 3 : reduce the graph by P4.



- Step 4 : reduce the graph by P3.



PROBLEM 3

- Consider the following snapshot of a system:

	<u>Allocation</u>	<u>Max</u>	<u>Available</u>
	<i>A B C D</i>	<i>A B C D</i>	<i>A B C D</i>
P_0	0 0 1 2	0 0 1 2	1 5 2 0
P_1	1 0 0 0	1 7 5 0	
P_2	1 3 5 4	2 3 5 6	
P_3	0 6 3 2	0 6 5 2	
P_4	0 0 1 4	0 6 5 6	

- Answer the following questions using the banker's algorithm:
 - a. What is the content of the matrix **Need**?
 - b. Is the system in a safe state?
 - c. If a request from process P_1 arrives for (0,4,2,0), can the request be granted immediately?

SOLUTION 3

- a. The values of **Need** for processes P_0 through P_4 respectively are $(0,0, 0, 0)$, $(0, 7, 5, 0)$, $(1, 0, 0, 2)$, $(0, 0, 2, 0)$, and $(0, 6, 4, 2)$.
- b. The system is in a safe state? Yes. With **Available** being equal to $(1, 5, 2, 0)$, either process P_0 or P_3 could run. Once process P_3 runs, it releases its resources, which allow all other existing processes to run.
- c. The request can be granted immediately? This results in the value of **Available** being $(1, 1, 0, 0)$. One ordering of processes that can finish is P_0, P_2, P_3, P_1 , and P_4 .